

VALGRIND

Linux x86 memory debugger and performance profiler

Aidan Shribman
SAP Labs Israel
November 2013



Agenda

Introduction

Detecting memory errors

Heap profiling and leak detection

Data concurrency race conditions

Profiling applications


Summary

What is Valgrind?

- **An open source system memory debugger:** Initiated by Julian Sward on 2000; released for Linux x86 on 2002 and winner of Open Source Award of 2004.
- **Simple and easy to use:** does not require re-compilation or re-linking: performs dynamic runtime instrumentation running on a synthetic CPU.
- **Used to validate many large Projects:** KDE, Emacs, Mozilla, OpenOffice and ... SAP HANA.
- **Programing language agnostic:** used for C/C++, Python, Java, JavaScript development.

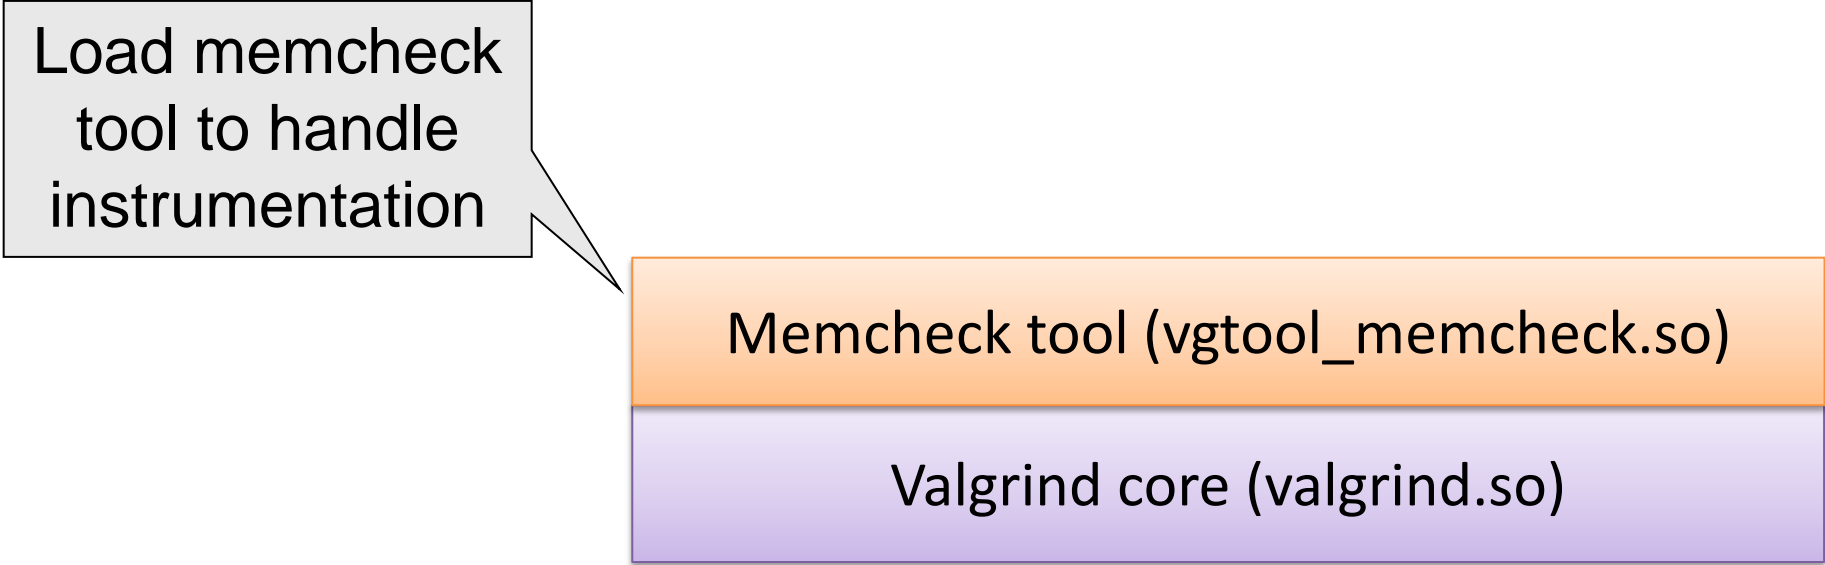
```
$ valgrind --tool=memcheck \  
    prog <args>
```

Uses
LD_PRELOAD
to load first

A diagram consisting of two rectangular boxes. The box on the left is light gray and contains the text 'Uses LD_PRELOAD to load first'. The box on the right is light purple and contains the text 'Valgrind core (valgrind.so)'. A black line connects the right side of the gray box to the left side of the purple box, forming a speech bubble shape that points towards the purple box.

Valgrind core (valgrind.so)

Load memcheck
tool to handle
instrumentation



```
graph LR; A[Load memcheck tool to handle instrumentation] --> B[Memcheck tool (vgtool_memcheck.so)]; B --- C[Valgrind core (valgrind.so)];
```

The diagram illustrates the process of loading a memcheck tool into the Valgrind core. A grey callout box on the left contains the text 'Load memcheck tool to handle instrumentation'. An arrow points from this box to an orange box labeled 'Memcheck tool (vgtool_memcheck.so)'. This orange box is stacked on top of a purple box labeled 'Valgrind core (valgrind.so)'.

Memcheck tool (vgtool_memcheck.so)

Valgrind core (valgrind.so)

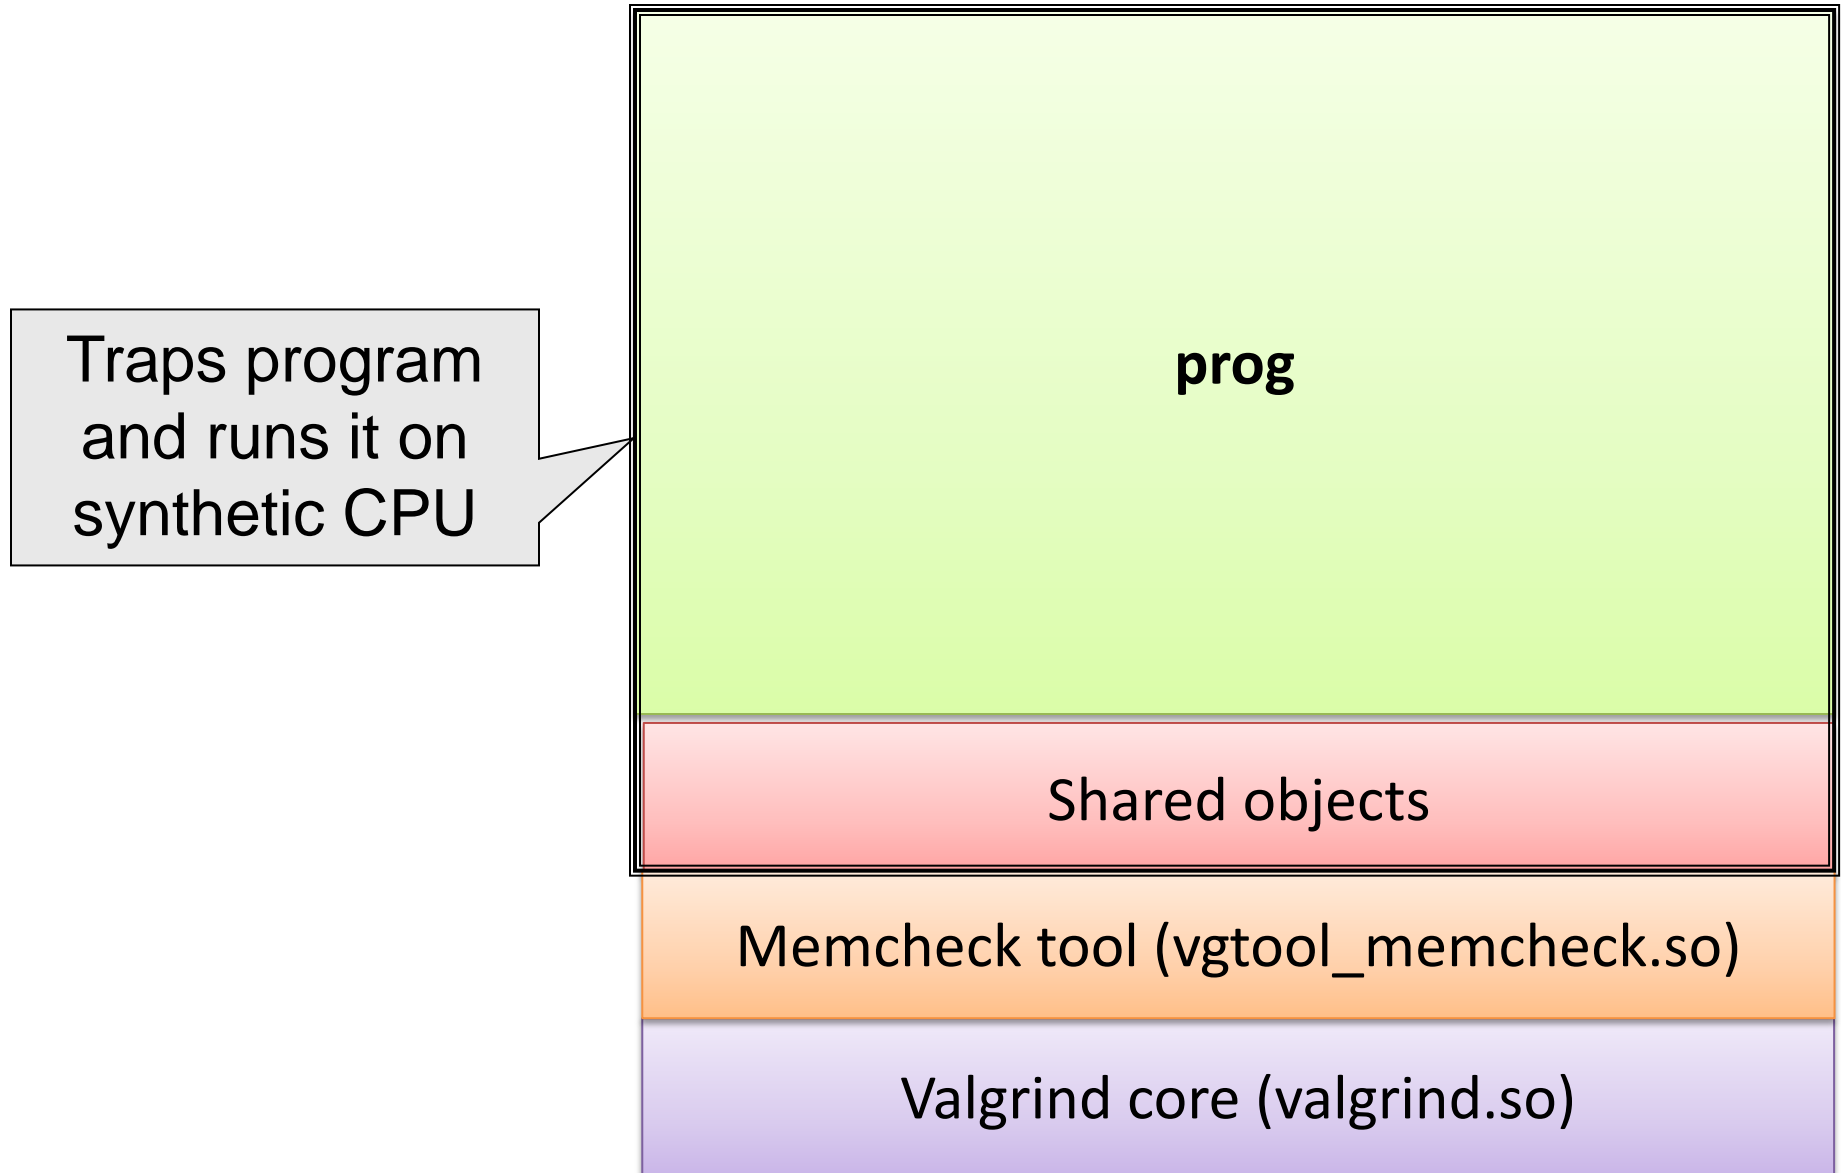
Traps program
and runs it on
synthetic CPU

prog

Shared objects

Memcheck tool (vgtool_memcheck.so)

Valgrind core (valgrind.so)



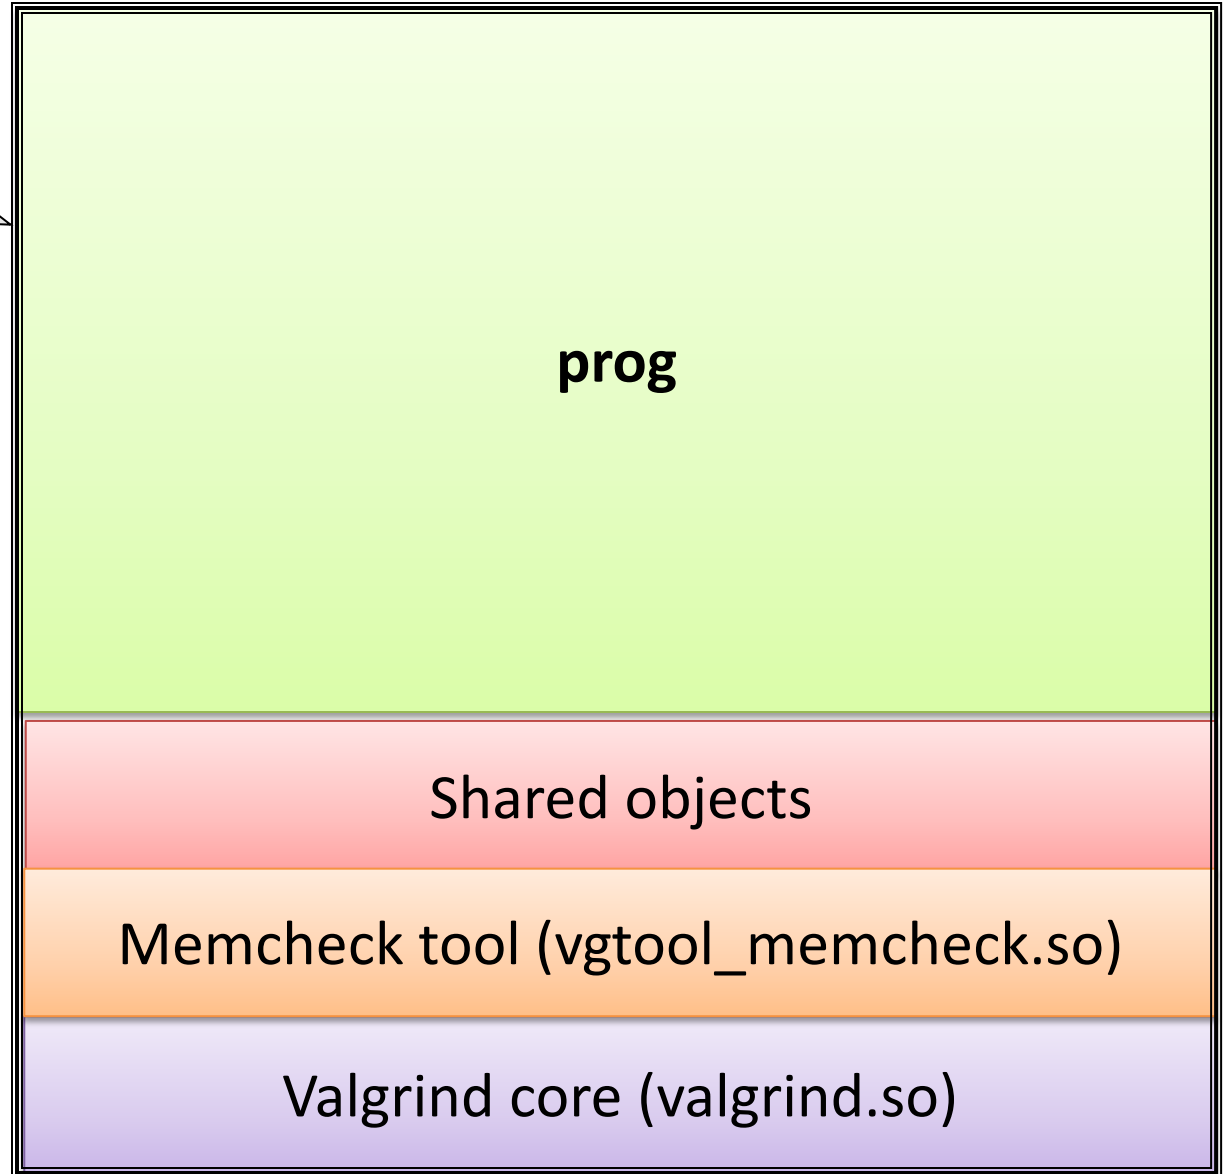
Appears as a
normal
application

prog

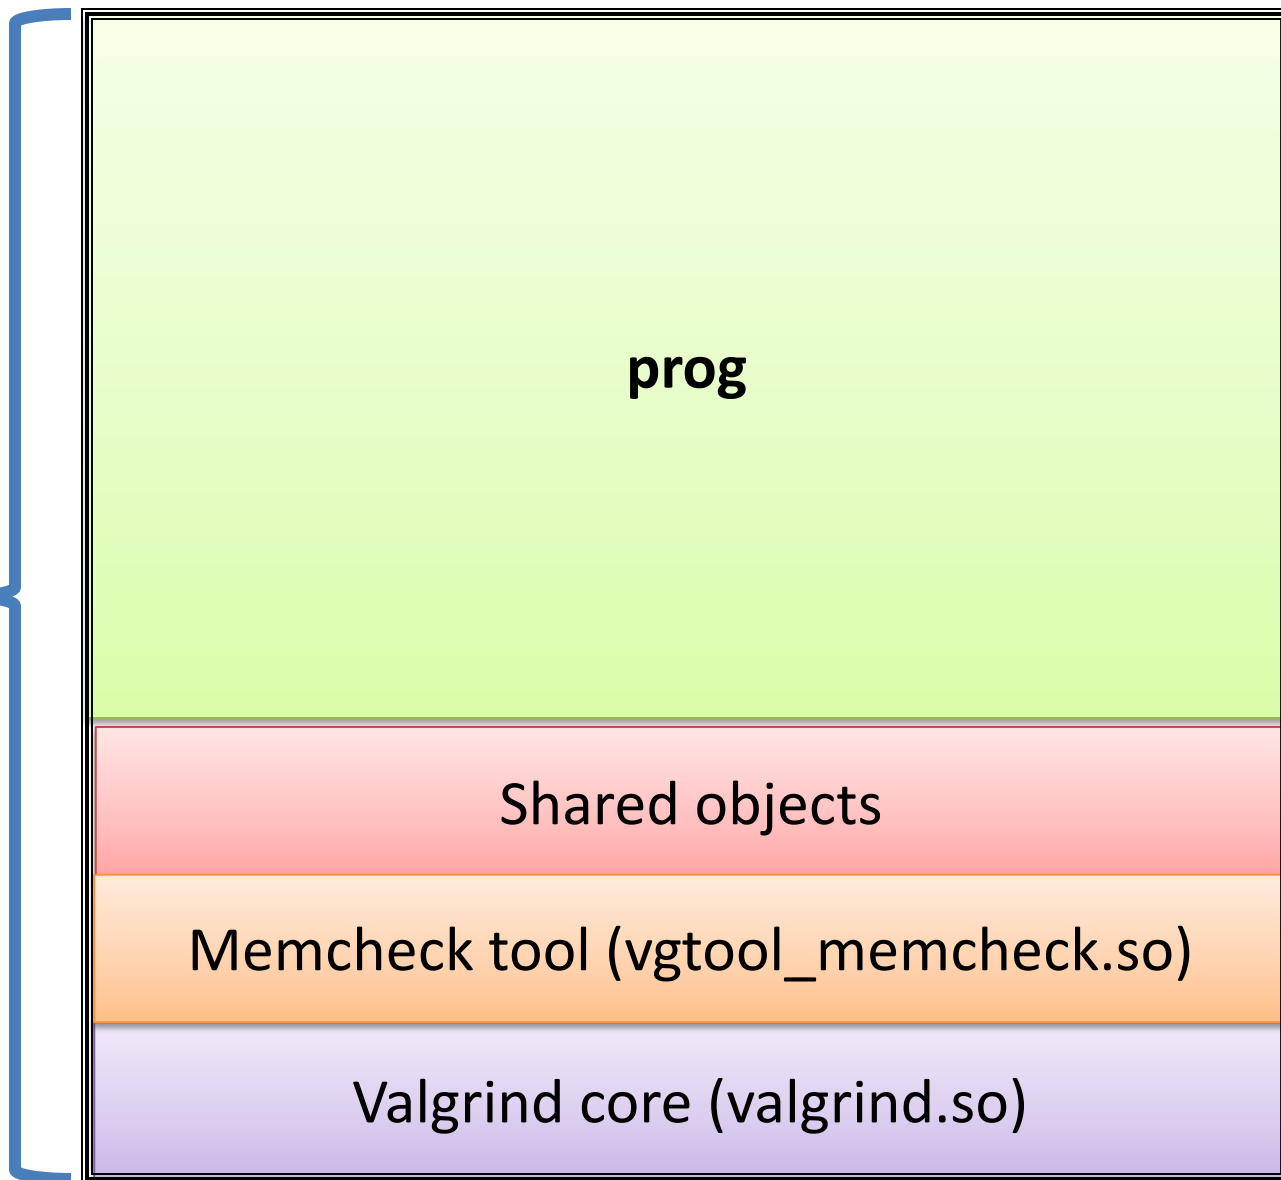
Shared objects

Memcheck tool (vgtool_memcheck.so)

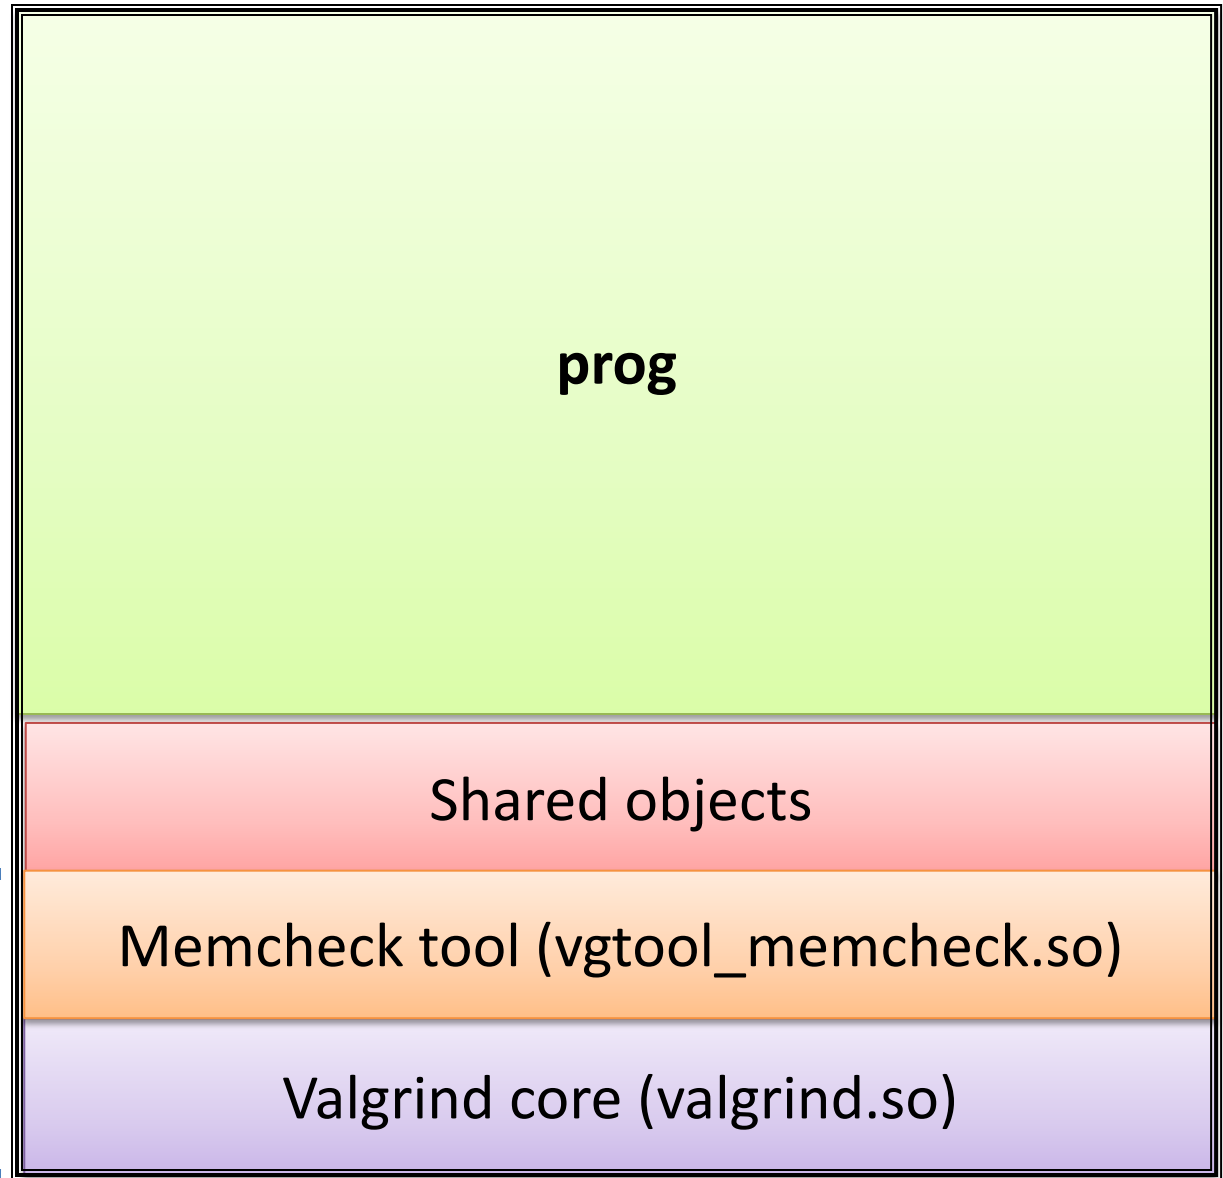
Valgrind core (valgrind.so)



Single Name Space



**PC (Program
Counter) under
Valgrind**



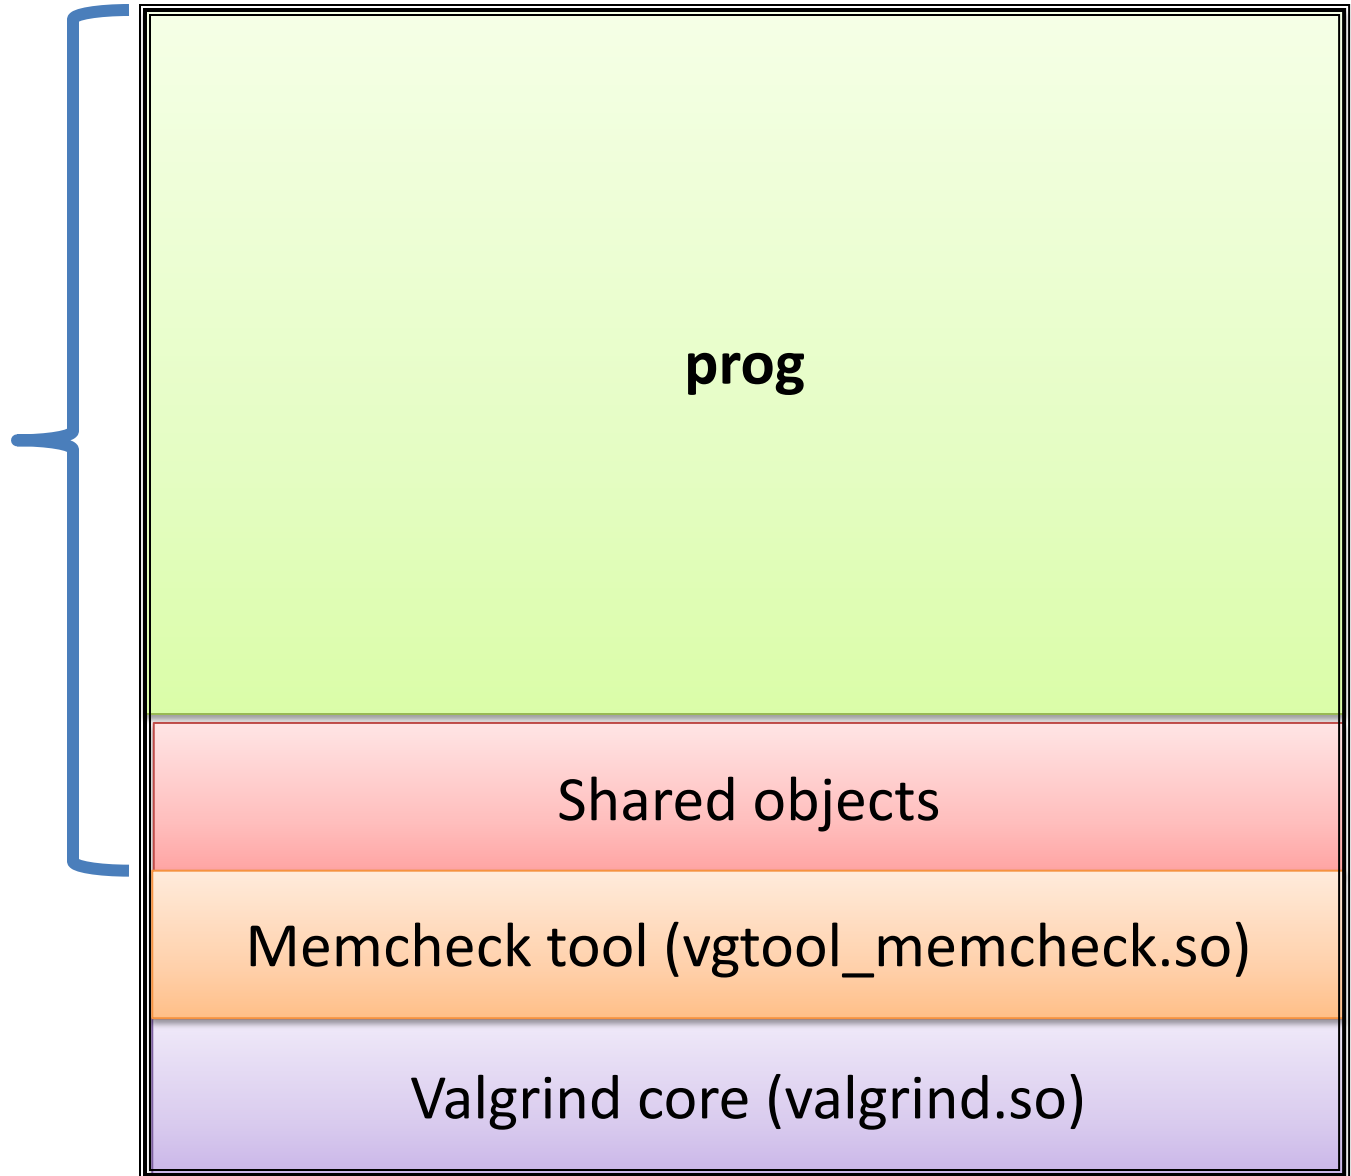
prog

Shared objects

Memcheck tool (vgtool_memcheck.so)

Valgrind core (valgrind.so)

**PC (Program
Counter) when
exists Valgrind**



Dynamic Binary Instrumentation

"UCode lies at the heart of the x86-to-x86 JITter. The basic premise is that dealing the the x86 instruction set head-on is just too darn complicated, so we do the traditional compiler-writer's trick and translate it into a simpler, easier-to-deal-with form."

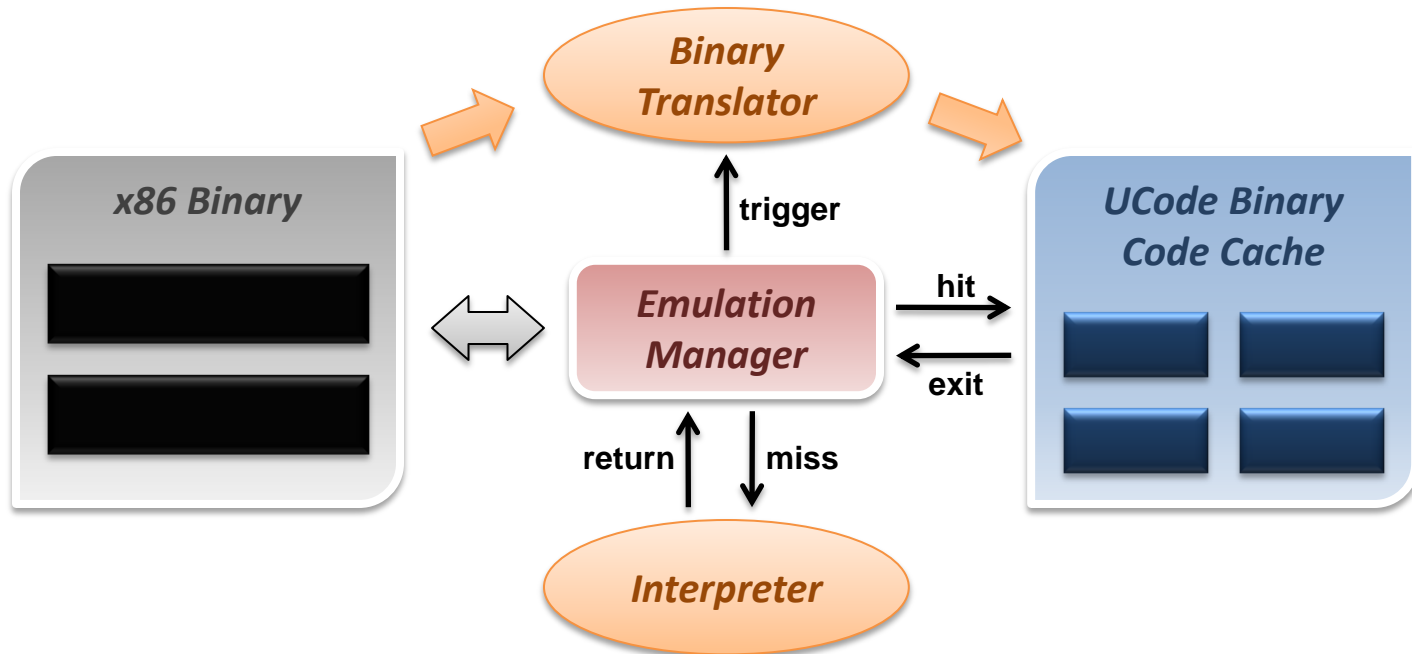
Julian Sward

The Instrumenting JITer

- Parse of an x86 basic block into a sequence of UCode instructions.
- UCode optimization with the aim of caching simulated registers in real registers.
- UCode instrumentation which adds value and address checking code.
- Post-instrumentation cleanup, removing redundant value-check computations.
- Register allocation done on UCode.
- Emission of final instrumented x86 code.

Dynamic Binary Translator

1. First time execution, no translated code in code cache.
2. Miss code cache matching, then directly interpret the guest instruction.
3. As a code block discovered, trigger the binary translation module.
4. Translate guest code block to host binary, and place it in the code cache.
5. Next time execution, run the translated code block from the code cache.



Installing Valgrind

```
$ zypper install valgrind
```

```
$ yum install valgrind
```

```
$ apt-get install valgrind
```

```
$ zypper install -y linux-kernel-headers
```

```
$ wget http://www.valgrind.org/downloads/valgrind-3.9.0.tar.bz2
```

```
$ tar xvfj valgrind-3.9.0.tar.bz2
```

```
$ cd valgrind-3.9.0
```

```
$ ./configure
```

```
$ make
```

```
$ make install
```

Memcheck

Memcheck

- Use of uninitialized memory
- Reading/writing memory after it has been freed
- Reading/writing off the end of malloc'd blocks
- Reading/writing inappropriate areas on the stack
- Memory leaks -- where pointers to malloc'd blocks are lost forever
- Passing of uninitialized or unaddressible memory to system calls
- Mismatched use of malloc/new/new[] vs free/delete/delete[]
- Some misuses of the POSIX pthreads API

- **A-bits**: every memory byte is shadowed by a single A (Addressability) bit indicating if the application can legitimately access the byte.
- **V-bits**: every memory byte is shadowed by eight V (Validity) bits indicating if the values of each bit are defined.
- **Heap blocks**: records heap blocks in auxiliary hash, enabling to track double-free and other miss-use.

```
$ valgrind --tool=memcheck \  
  --num-callers=20 \  
  --log-file=vg.log \  
  --trace-children=yes \  
  prog <args>
```

```
int v, *p = malloc(10 * sizeof(int));  
v = p[10]; /* invalid read (heap) */
```



```
int a[10];
```

```
a[10] = 1; /* invalid write (stack) */
```

Not Detected!

```
int a, b;  
b = a; /* pass uninitialized */  
if (!b) /* use uninitialized */  
    printf("non deterministic\n");
```

```
char *p = malloc(100); /* uninitialized */  
write(1, p, 100); /* to syscall */
```

```
unsigned char c;  
c |= 0x1; /* only 1st bit is initialized */  
if (c & 0x1) /* bit-1 initialized */  
    printf("always be printed\n");  
if (c & 0x2) /* bit-2 uninitialized */  
    printf("non deterministic\n");
```

Addrcheck

Addrcheck

- ~~Use of uninitialized memory~~
- Reading/writing memory after it has been freed
- Reading/writing off the end of malloc'd blocks
- Reading/writing inappropriate areas on the stack
- Memory leaks -- where pointers to malloc'd blocks are lost forever
- Passing of ~~uninitialized~~ or unaddressible memory to system calls
- Mismatched use of malloc/new/new[] vs free/delete/delete[]
- Some misuses of the POSIX pthreads API

```
$ valgrind --tool=addrcheck \  
    prog <args>
```

Performance slowdown

Program	Time (s)	Nulgrind	Memcheck	Addrcheck	Cachegrind
bzip2	10.7	2.4	13.6	9.1	31.0
crafty	3.5	7.2	44.6	26.5	107.4
gap	0.9	5.4	28.7	14.4	46.6
gcc	1.5	8.5	36.2	23.6	73.2
gzip	1.8	4.4	20.8	14.5	50.3
mcf	0.3	2.1	11.6	5.9	18.5
parser	3.3	3.7	17.4	12.5	34.8
twolf	0.2	5.2	29.2	18.5	53.3
vortex	6.5	7.5	47.9	32.7	88.4
ammp	18.9	1.8	24.8	21.1	47.1
art	26.1	5.9	14.1	11.5	19.4
equake	2.1	5.5	32.7	28.0	49.9
mesa	2.7	4.7	41.9	31.6	64.5
median		5.2	28.7	18.5	49.9

Memory increase

Program	Size (KB)	Nulgrind	Memcheck	Addrcheck	Cachegrind
bzip2	34	5.2	12.1	6.8	9.1
crafty	156	4.5	10.9	5.9	8.2
gap	140	5.6	12.7	7.3	9.7
gcc	564	5.9	13.1	7.6	9.9
gzip	30	5.5	12.6	7.2	9.4
mcf	30	5.7	13.5	7.7	9.9
parser	97	6.0	13.6	7.8	10.1
twolf	114	5.2	12.2	7.0	9.3
vortex	234	5.8	13.2	8.1	10.1
ammp	68	4.7	11.7	7.1	9.5
art	24	5.5	13.0	7.5	9.8
equake	44	5.0	12.2	7.1	9.2
mesa	69	4.8	11.2	6.7	8.9
median		5.5	12.6	7.2	9.5

Getting Accurate Backtraces

gcc main.c && strip a.out

==23580== Invalid write of size 1

--23580== at 0x8048622: (within /root/dev/demo/test3)

==23580== by 0x8048703: (within /root/dev/demo/test3)

==23580== by 0x42015573: __libc_start_main (in /lib/tls/libc-2.3.2.so)

==23580== by 0x80484A0: (within /root/dev/demo/test3)

==23580== Address 0x3C03502E is 0 bytes after a block of size 10 alloc'd

==23580== at 0x3C01E250: malloc (vg_replace_malloc.c:105)

==23580== by 0x8048615: (within /root/dev/demo/test3)

==23580== by 0x8048703: (within /root/dev/demo/test3)

==23580== by 0x42015573: __libc_start_main (in /lib/tls/libc-2.3.2.so)

gcc main.c

==23587== Invalid write of size 1

--23587== at 0x8048622: invalid_write (in /root/dev/demo/test3)

==23587== by 0x8048703: main (in /root/dev/demo/test3)

==23587== Address 0x3C03502E is 0 bytes after a block of size 10 alloc'd

==23587== at 0x3C01E250: malloc (vg_replace_malloc.c:105)

==23587== by 0x8048615: invalid_write (in /root/dev/demo/test3)

==23587== by 0x8048703: main (in /root/dev/demo/test3)

gcc -g main.c

==23594== Invalid write of size 1

--23594== at 0x8048622: invalid_write (demo_memcheck.c:34)

==23594== by 0x8048703: main (demo_memcheck.c:68)

==23594== Address 0x3C03502E is 0 bytes after a block of size 10 alloc'd

==23594== at 0x3C01E250: malloc (vg_replace_malloc.c:105)

==23594== by 0x8048615: invalid_write (demo_memcheck.c:31)

==23594== by 0x8048703: main (demo_memcheck.c:68)

Error Suppression

```
$ valgrind --tool=memcheck \  
    --gen-suppressions=yes \  
    --logfile=my.supp \  
    prog <args>
```

{

<gdk_set_locale>

Memcheck:Cond

...

fun:gdk_set_locale

}


```
{
```

```
<libpango>
```

```
Memcheck:Leak
```

```
...
```

```
obj:/usr/*lib*/libpango*
```

```
}
```

```
$ valgrind --tool=memcheck \  
    --suppressions=my.supp.* \  
    prog <args>
```

Attaching a Debugger

```
#include <malloc.h>
```

```
int main()
```

```
{
```

```
    int i;
```

```
    char *a = malloc(3);
```

```
    for (i=0; i<4; i++)
```

```
        a[i] = 0;
```

```
}
```

```
$ valgrind --tool=memcheck \  
    --db-attach=yes \  
    ./prog <args>
```

```
==3513== Memcheck, a memory error detector
==3513== Copyright (C) 2002-2010, and GNU GPL'd, by Julian Seward et al.
==3513== Using Valgrind-3.6.0.SVN-Debian and LibVEX;
==3513== Command: ./a.out
==3513==
==3513== Invalid write of size 1
==3513== at 0x40051C: main (in /home/ortyl/a.out)
==3513== Address 0x51b1043 is 0 bytes after a block of size 3 alloc'd
==3513== at 0x4C2815C: malloc (vg_replace_malloc.c:236)
==3513== by 0x400505: main (in /home/ortyl/a.out)
==3513==
==3513==
==3513== ---- Attach to debugger ? --- [Return/N/n/Y/y/C/c] ---- y
==3513== starting debugger with cmd: /usr/bin/gdb -nw /proc/3516/fd/1014 3516
(gdb)
```

Leak Check

```
$ valgrind --tool=memcheck \  
  --leak-check=yes \  
  --show-reachable=yes \  
  prog <args>
```


HEAP SUMMARY:

in use at exit: 4 bytes in 1 blocks
total heap usage: 1 allocs, 0 frees, 4 bytes allocated

4 bytes in 1 blocks are still reachable in loss record 1 of 1
at 0x4024C1C: malloc (vg_replace_malloc.c:195)
by 0x40B0CDF: strdup (strdup.c:43)
by 0x804879B: main (in /home/aleksander/valgrind-memcheck)

LEAK SUMMARY:

definitely lost: 0 bytes in 0 blocks
indirectly lost: 0 bytes in 0 blocks
possibly lost: 0 bytes in 0 blocks
still reachable: 4 bytes in 1 blocks
suppressed: 0 bytes in 0 blocks

```
void *rrr;
```


```
int main(void)
```

```
{
```

```
    rrr = strdup("bbb") ;
```

```
    return 0 ;
```

```
}
```



Reachable

```
void *rrr;
```

```
int main(void)
```


```
{
```

```
    rrr = strdup("bbb") ;
```

```
    rrr = NULL;
```


```
    return 0 ;
```

```
}
```



Defenetely Lost

Directly Lost



RRR

0x0

AAA

0x9d90FF0

BBB

'b'

'b'

'b'

\0



Indirectly Lost



```
void *rrr ;
```

```
int main(void)
```

```
{
```


```
    rrr = strdup("bbb") ;
```

```
    rrr = ((char *) rrr) + 1;
```

```
    return 0 ;
```

```
}
```

Possibly Lost



Possibly Lost

RRR

0x9d90FF1

AAA

0x9d90009

BBB

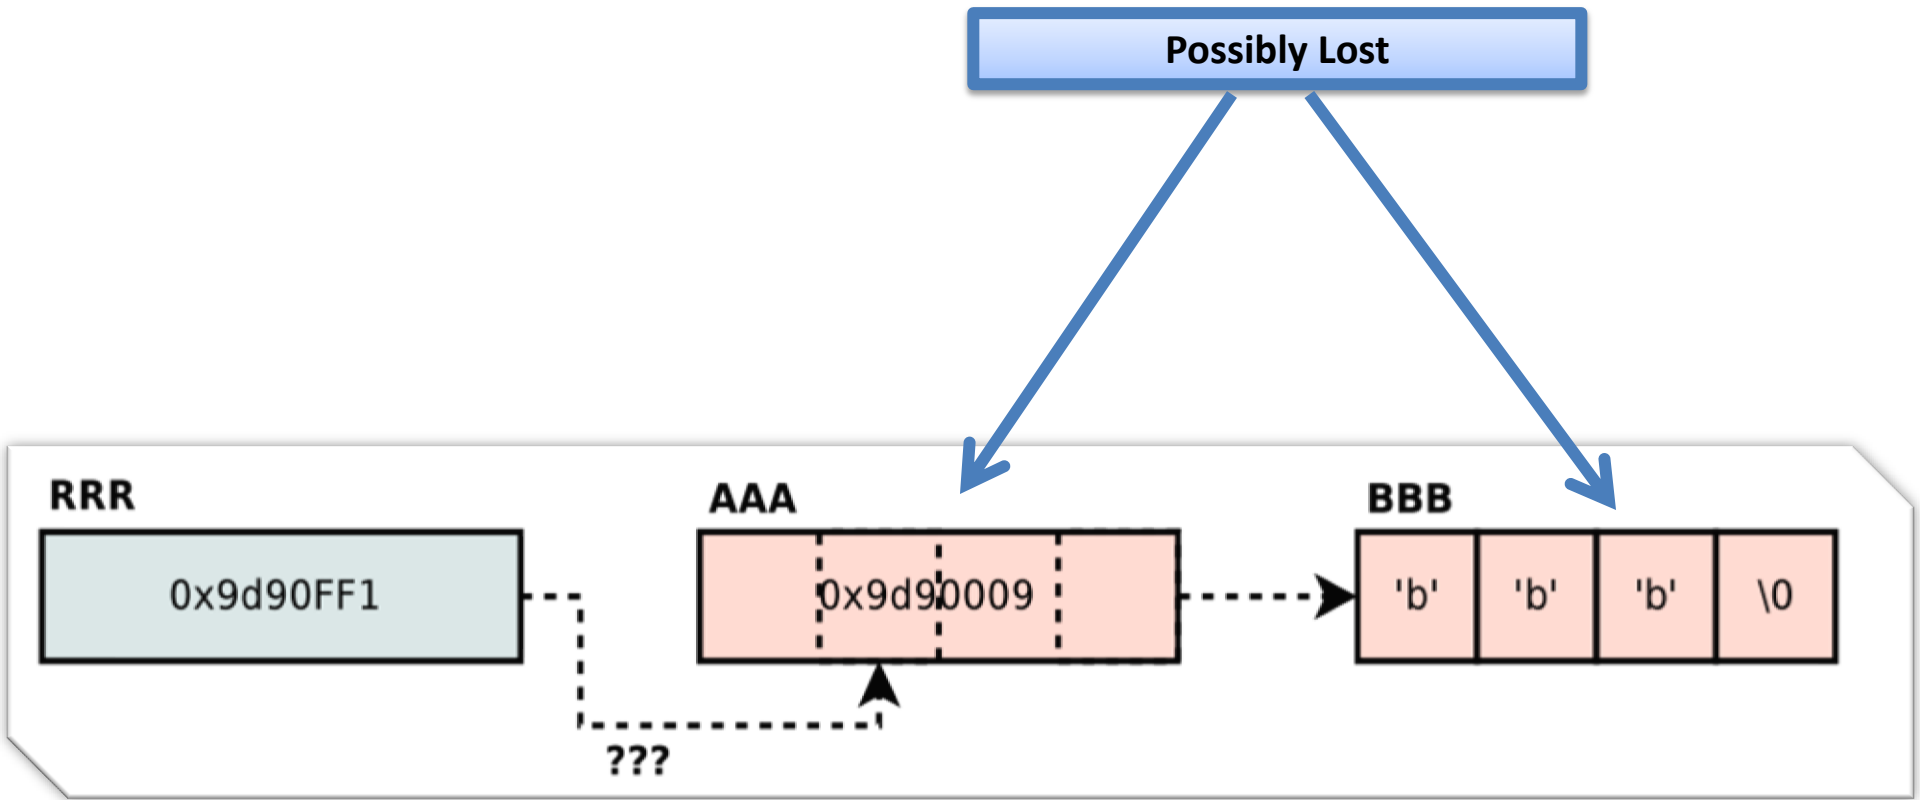
'b'

'b'

'b'

\0

???



```
$ valgrind --tool=memcheck \  
    --vgdb-error=0 --vgdb=yes \  
    prog <args>
```

(gdb) target remote | vgdb

(gdb) monitor leak_check full reachable any

==2418== 100 bytes in 1 blocks are still reachable in loss record 1 of 1

==2418== at 0x4006E9E: malloc (vg_replace_malloc.c:236)

==2418== by 0x804884F: main (prog.c:88)

==2418==

==2418== LEAK SUMMARY:

==2418== definitely lost: 0 bytes in 0 blocks

==2418== indirectly lost: 0 bytes in 0 blocks

==2418== possibly lost: 0 bytes in 0 blocks

==2418== still reachable: 100 bytes in 1 blocks

==2418== suppressed: 0 bytes in 0 blocks


==2418==

(gdb)

Massif

Massif is a heap profiler: it measures how much heap memory your program uses. This includes both the useful space, and the extra bytes allocated for book-keeping and alignment purposes. It can also measure the size of your program's stack(s), although it does not do so by default.

Memory consumption of bin/asf_harness
Peak of 2,7 MiB at snapshot #55



Helgrind

Helgrind is a thread debugger: finds data races in multithreaded programs. It looks for memory locations which are accessed by more than one (POSIX p-) thread, but for which no consistently used (pthread_mutex_) lock can be found.


```
static void *inc_shared (void *v) {  
    shared++; /* un-protected access to shared */  
    return 0;  
}
```

```
pthread_t a, b;  
pthread_create(&a, NULL, inc_shared, NULL);  
pthread_create(&b, NULL, inc_shared, NULL);
```

Callgrind

Callgrind is a profiling tool: it records the call history among functions in a program's run as a call-graph. Collecting number of instructions executed, their relationship to source lines, the caller/callee relationship between functions, and the numbers of such calls.


```
$ valgrind --tool=callgrind prog <args>
```

```
$ callgrind_control --zero # zero counters
```

... the application runs ...

```
$ callgrind_control --dump # dump counters
```

```
$ kcache-grind cachegrind.out.<pid>
```



Parts | Types | Callers | Source

The screenshot shows the Windows Task Manager Performance tab. The 'Processes' section is expanded, showing a list of processes with their names and CPU usage percentages. The processes listed are:

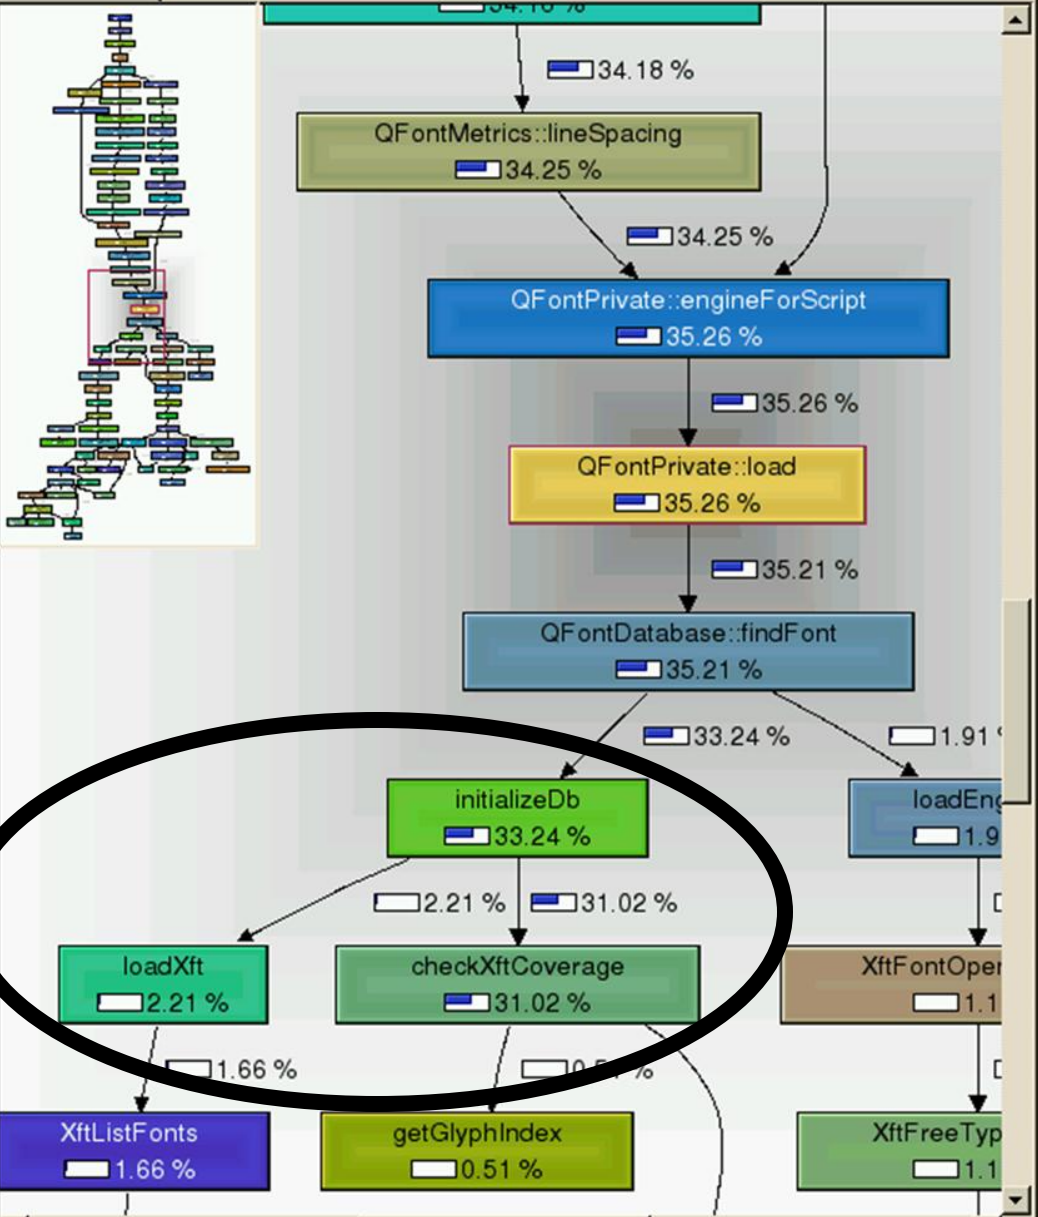
- T1_Get_Private_Dict: 9.47%
- T1_Decrypt: 3.34%
- hexa_value: 2.31%
- parse_dict: 11.49%
- PS_Unicode_Value: 5.51%

The 'Performance' section shows the following metrics:

- Free Space: 96%
- Free Memory: 100%

The 'System' section shows the following metrics:

- Free Space: 96%
- Free Memory: 100%



Profiling golden rules

- A human has very seldom a clue at all where in the code time is spent, therefore you must use a profiler.
- You should never optimize code unless the application feels slow as optimization always leads to code that is harder to maintain.
- Make sure to profile your application in a realistic context.

Eclipse



Create, manage, and run configurations

Profile C/C++ Application Using Valgrind



type filter text

▼ Profile With Valgrind

▼ alloctest Debug

Name: alloctest Debug

Main Arguments Valgrind Options Environment Source Common

Tool to run: memcheck ▼

General Tool

Basic Options

trace children on exec: ☐child silent after fork: ☒run __libc_freeres() on exit: ☒

Error Options

demangle C++ names: ☒

num callers in stack trace: 12 ▲▼

limit errors reported: ☒show errors below main: ☐

max size of stack frame: 2000000 ▲▼

suppressions file:

Workspace...

File System...

Apply

Revert

Profile

Close

Filter matched 2 of 3 items



File Edit Refactor Navigate Search Run Project Window Help



C/C++

Project Explorer

- alloctest
- exectest
- External Plug-in Libraries
- forktest
- randalloctest
- test
 - Binaries
 - Includes
 - Debug
 - test.c
 - test.txt
 - test supfile.supp

forktest.c alloctest.c test.c

```
/*
 * test.c
 *
 * Created on: Sep 12, 2008
 * Author: ebaron
 */
#include <stdlib.h>
#include <stdio.h>
#include <fcntl.h>

#define SIZE 10
int main() {
    open("../test.c", O_RDONLY);
    char *waste = (char *)malloc(sizeof(char) * SIZE);
    int *a;
    printf("%d\n", *a);
    fgets(waste, SIZE, stdin);
    waste[SIZE] = 0;
    return 0;
}
```

Outline Make Tar

- stdlib.h
- stdio.h
- fcntl.h
- # SIZE
- main() : int

Problems Tasks Console Properties Valgrind

test Debug [memcheck] valgrind (12/9/08 4:03 PM)

- Use of uninitialised value of size 4 [pid: 16002 / tid: 1]
 - at 0x80484AB: main (test.c:16)
- Invalid write of size 1 [pid: 16002 / tid: 1]
- 10 bytes in 1 blocks are definitely lost in loss record 1 of 1 [pid: 16002 / tid: 1]


File Edit Refactor Navigate Search Run Project Window Help



Project Explorer

- ▶ allocetest
- ▶ exectest
- ▶ External Plug-in Libraries
- ▶ forktest
- ▼ gnome-terminal
 - ▼ Binaries
 - ▶ gnome-terminal - [x86/le]
 - ▶ Archives
 - ▶ Includes
 - ▶ build
 - ▶ src
- ▶ randalloctest
- ▶ test

Heap Chart - gnome-terminal



Outline

Make Tar

An outline is not available.

Problems Tasks Console Properties Valgrind

gnome-terminal [massif] valgrind (12/22/08 10:29 AM)

- ▶ Snapshot 59 - 92.75% (1373487B) (heap allocation functions) malloc/new/new[], --alloc-fns, etc.
- ▶ Snapshot 62 - 92.59% (1387525B) (heap allocation functions) malloc/new/new[], --alloc-fns, etc.
- ▼ Snapshot 65 - 92.39% (1399350B) (heap allocation functions) malloc/new/new[], --alloc-fns, etc.
 - ▼ 69.23% (1048576B) 0x4698AF8: google_breakpad::MinidumpGenerator::AllocateStack() (in /usr/lib/gtk-2.0/modules/libgnomebreakpad.so)
 - ▼ 69.23% (1048576B) 0x4698BAA: google_breakpad::MinidumpGenerator::MinidumpGenerator() (in /usr/lib/gtk-2.0/modules/libgnomebreakpad.so)
 - ▼ 69.23% (1048576B) 0x469760E: google_breakpad::ExceptionHandler::ExceptionHandler(std::string const&, bool (*)(void*), bool (*)(char const*, char const*, void*, bool))
 - ▼ 69.23% (1048576B) 0x469684D: gtk_module_init (in /usr/lib/gtk-2.0/modules/libgnomebreakpad.so)
 - ▼ 69.23% (1048576B) 0x25ED6DC: (within /usr/lib/libgtk-x11-2.0.so.0.1400.5)

Language Support

Language Survey

- **C: 56**
- **C+: 52**
- Fortran: 6
- **Java: 3**
- asm: 3
- **Python: 2**
- TCL/TK: 1
- **Objective C: 1**
- Others: 3

Python

```
$ cd python/dist/src
```

```
$ valgrind --tool=memcheck \  
  --suppressions=Misc/valgrind-python.supp \  
  ./python -E -tt ./Lib/test/regrtest.py -u \  
  bsddb,network
```

Node.js

```
$ valgrind --leak-check=yes \  
node foo.js
```

Java JNI

```
$ valgrind --trace-children=yes \  
  --leak-check=full \  
  java -Djava.library.path=$(PWD) Foo
```

one slide to go ... 😊

- Valgrind is a robust x86-Linux memory debugger.
- Using it regularly enhances product's quality.
- When to use it?
 - All the time.
 - In automatic testing.
 - After big changes.
 - When a bug occurs.
 - When a bug is suspected.
 - Before a release.

THANK YOU!

Valgrind

www.valgrind.org

Aidan Shribman;

SAP Labs Israel

aidan.shribman@sap.com



References

- Aleksander Morgado. *Understanding Valgrind memory leak reports*. February 4, 2010
- Nicholas Nethercote and Julian Seward. *How to Shadow Every Byte of Memory Used by a Program*. 2007.
- Nicholas Nethercote and Julian Seward. *Valgrind: A Program Supervision Framework*. 2003.